

Honours Dissertation: Semi-Supervised Learning on Data Streams

Michael Glenny

October 28, 2014

0.1 Abstract

Introduction

Recent developments in data collection techniques have now left researchers and businesses drowning in data. This has meant that data mining often cannot be performed in the static manner in which it once was. Instead, we need novel techniques to handle not only the incredible volume of data arriving, but also the fact that the nature of this data can change over time. Previous data stream classification techniques frequently relied on entirely labelled data [9, 10]. However, due to the volume of current data streams, completely labelling a data stream may be simply infeasible. For instance, the amount of transactions going through a bank are far too voluminous for even a large team of experts to manually classify every single transaction. In response, a new collection of techniques have arisen which focus on labelling only a small part of the data stream, and using this limited information to classify the remainder of the stream [1, 2].

Sensitive to these concerns, in this paper we will attempt to improve the classification accuracy of Label Stream [1], a current semi-supervised approach by scrutinising the quality of the clustering it performs. My main contribution is to attempt to use outlier detection measures to evaluate or change the clusters produced by the semi-supervised clustering technique in an attempt to improve the accuracy of the classification. We will evaluate the quality of the clustering performed by a model using LoOP, the Local Outlier Probability score [13], and attempt to improve accuracy by either pruning outliers or reducing the influence of poor quality clusters.

The structure of the paper is as follows. In the following section, we will discuss the various ways researchers have confronted the problem of classifying data streams with limited labelled data. In Section 3, we will review how a model is built on a partially labelled data chunk in the Label Stream algorithm, and in Section 4 we will investigate the function of the ensemble in the Label Stream algorithm. Then in Section 5 we will investigate the Local Outlier Probability and how we can use it to evaluate clusters. In Section 6 we will review experimental results on several data sets. Finally in Section 7 we will conclude and discuss future work.

Related Work

There have been several different approaches to handling the problem of semi-supervised data stream classification, and many of the main approaches have been adapted to attempt to handle the problems of sparse labelled data.

One of these approaches is the REDLLA algorithm, developed by Li et al. [11]. This algorithm adapts the CVFDT (Concept-adapting Very Fast Decision Tree) [8], a decision tree approach to classifying a data stream, to handle semi-labelled data streams. One of the strengths of this approach is that because a decision tree can store a great amount of classification information, REDLLA is especially adept at handling recurring concept drift. However, the trade off for this benefit is that the decision tree can occupy a lot of space in memory, as noted by in the limitations section by the authors.

Another approach is was developed by Zhang et al. [2], which involves a slightly different framework. Here, relational k-mean clustering is combined with a Support Vector Machine algorithm to create a semi-supervised technique which copes admirably with a lack of labels in the data stream. This algorithm greatly out-performed previous approaches based on Support Vector Machines when concept drift and partially labelled streams were involved.

Finally, as we will see, the Label Stream algorithm developed by Woolam et al. [1] uses a semi-supervised K-means clustering technique on top of a label propagation algorithm to classify unlabelled clusters. Using an ensemble gives a natural way to detect concept drift without using too much memory, but is less appropriate in cases where there is recurring drift. However, it does a fixed amount of memory is devoted to tracking concept drift.

These approaches all have something in common; they all use clustering to associate unlabelled points with labelled points, and then they all apply different supervised techniques to leverage what little labelled data they possess for predictions on new unlabelled points.

/*Discussion of incremental vs ensemble*/

Label Stream: An Algorithm for Classifying Stream Data

The algorithm that I chose to attempt to improve was one developed by researchers at the University of Dallas. The Label Stream algorithm was pub-

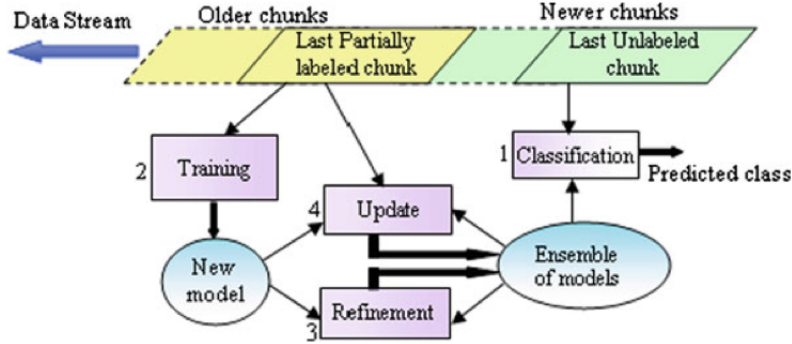


Figure 1: An overview of the flow of control in the Label Stream algorithm [3]

lished in 2009 by Clay Woolam, Mohammad M. Masud and Latifur Khan [1]. I was drawn to this algorithm for several reasons. Firstly, an ensemble approach seems like an extremely natural and elegant way to approach the problem of detecting concept drift in streams. Secondly, I was rather curious about semi-supervised clustering, and thirdly, the fact that Label Stream worked with fixed memory also seemed an added bonus.

The basic framework of the algorithm is to use an ensemble of models to predict which class a data point belongs to. The stream is broken into chunks, and each model is built from the data in a partially labelled chunk. When a new model is created, we simply need to update our ensemble by evicting the model which is performing worst on the new labelled data, as well as refining the current models in the ensemble based on the labelled data seen in the new chunk.

Training a Model

We can train a model from a data chunk as soon as some arbitrary percentage of the data points in that chunk have been manually labelled. Experimental testing by Woolam et al. shows that Label Stream performs as well as many previous stream classification algorithms with as few as 10% labelled points [1].

Training a model based on a chunk of data consists of three steps. Firstly, we cluster the data into large macro-clusters based on a modified version of K-

Means clustering, in which we try not only to minimise intra-cluster distance like in standard K-means, but also minimise cluster impurity by taking into account the entropy of clusters as a dissimilarity measure in which clusters with points of different labels are punished.

Secondly, we then split those macro-clusters into class label pure micro-clusters. Some of these micro-clusters will be unlabelled, and hence the third step is to use a Label Propagation technique to label the unlabelled micro-clusters. Now that all the micro-clusters are labelled, they are ready to be used in the classification process.

The Macro-Clustering Step

MCI K-Means

In regular K-means clusters, we wish to minimise the intra-cluster distance in every cluster, which can be expressed as minimising an objective function:

$$O_{Dist} = \sum_{i=1}^K \sum_{x \in X_i} ||x - \mu_i||^2$$

where K is the number of clusters, X_i is cluster i , x is a data point and μ_i is the centroid of cluster i . This is an unsupervised technique, but we would like to take advantage of the fact that a certain percentage of our data chunk have labels associated. To do this, we will modify unsupervised K-means by adding an additional factor to the objective function. This is called by Woolam et al. MCI K-means, which stands for Minimising Cluster Impurity for K-means [1]. The objective function for this variation on K-means is:

$$O_{MCIKMeans} = \sum_{i=1}^K \sum_{x \in X_i} ||x - \mu_i||^2 (1 + Ent_i ADC_i)$$

This is the same as the previous function except for Ent_i , the entropy of cluster i and ADC_i is the aggregated dissimilarity count. However, to understand this formula, we need to define these terms. Entropy is a measure which is normally used to describe how many bits are necessary to convey information, but is useful in this context as a measure of purity of clusters. The formula for calculating the entropy of a cluster i is:

$$Ent_i = \sum_{c=0}^C (-p_i^c \log(p_i^c))$$

where p_i^c is the prior probability of a point in the cluster i being of class c , i.e.:

$$p_i^c = \frac{|X_i(c)|}{|X_i|}$$

where $|X_i(c)|$ is the number of points of class c in cluster i and $|X_i|$ is the number of points in cluster i . Entropy will give a higher penalty to impure clusters, however using only entropy will mean that every point has the same penalty related to impurity. This is a problem when we try and minimise our cluster impurity. Suppose for instance we have cluster X, with 9 points of label A, and cluster B, with 3 points of label A. Then, suppose we add 7 points of label B.

The best arrangement we can hope for is for cluster X to remained untouched, and cluster Y to take all the points of label B, resulting in an entropy of 0.26. However, what will happen when we incrementally add each point? Well, in the first instance, because cluster X is larger, a point of label B will be assigned to it. Now because cluster X is impure, we will continue adding points to cluster X, and our final configuration will be cluster X with 16 points in total, 9 of class A and 7 of class B, and an entropy of 0.29, and cluster Y as it was with an entropy of 0. This is not the optimal configuration!

To solve this problem, we introduce a measure of cluster impurity called Dissimilarity Count. We define the dissimilarity count of a point x in a cluster i to be:

$$DC(x, i) = \sum_{c=0}^C (|X_i| - |X_i c|)$$

Basically, the dissimilarity count of a single point is the number of points in the cluster with a class label other than the class of that point. In our previous example, using Dissimilarity Count, we would place all the points of label B into cluster Y, because the dissimilarity count for all points of label B would stay the constant 3, for all the points of label A which were there before.

The Aggregated Dissimilarity Count, ADC_i , is simply the sum of all dissimilarity counts for each point in a cluster i .

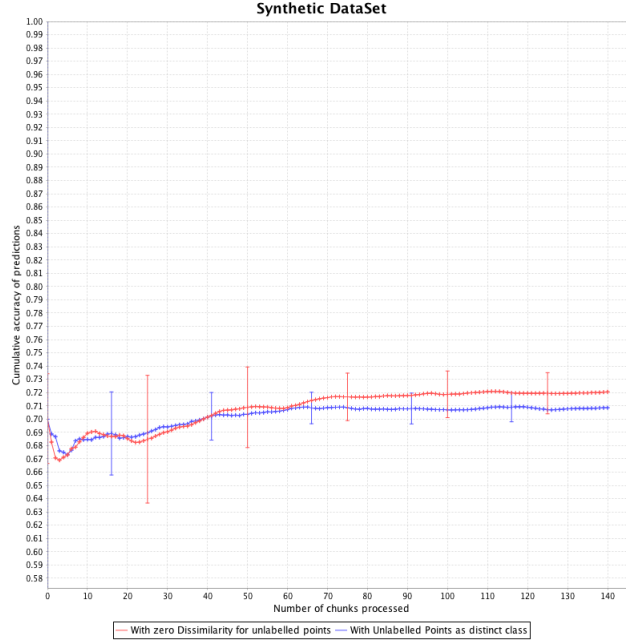


Figure 2: Considering unlabelled points as a distinct class for the purposes of dissimilarity count slightly improves classification accuracy.

An outstanding issue is whether to consider unlabelled points as a distinct class or not. In the original Label Stream paper, it is clear that unlabelled points are not to be considered in the Dissimilarity count, as they explicitly state that if a point is unlabelled, $DC(x, i) = 0$ [1]. However, in Khan et al. review of the same algorithm, it is more equivocal, and the paper seems to suggest that they did consider unlabelled points as a distinct class for the purposes of calculating the dissimilarity count [3]. For my experiments, I have decided to adopt the approach of not considering unlabelled points as distinct classes, and hence for any unlabelled point, $DC(x, i) = 0$. I did this because I found the Khan et al. paper somewhat equivocal on this matter and on experimental tests, this matter performed slightly better. We can see evidence of this in Figure 2.

Expectation Minimisation

In order to generate these clusters, firstly we need to initialise our clusters. This initialisation can have a great effect on the quality of the subsequent clustering. Firstly, we take all the labelled points in our data chunk and calculate the prior probability of a certain class appearing in that chunk.

Then, for each class, we calculate the proportion of the K clusters to be generated which should be initialised to a point of that class. This calculation is:

Number of clusters initialised to class $i = p_i^c K$

where K is the number of macro-clusters. Now, we use the furthest apart initialisation method to select the points which will be the centroids of our macro-clusters. That is, we find the point of class i which is furthest away from all the points in i which have already been selected as centroids for clusters. Masud et al. report that this method of initialisation, proportionate to the distribution of labelled data in the data chunk greatly improved their cluster quality [3].

With the clusters initialised, we can now apply Expectation Minimisation to find a clustering. Firstly, we apply the E-step, in which we assign a point to the cluster which minimises the contribution to the global function, regardless of whether that point is already assigned to a cluster or not:

$$O_{MCIKMeans}(x) = ||x - \mu_i||^2(1 + Ent_i DC_i(x))$$

The value of the overall objective function depends on the order by which we add points from the chunk to the clusters. If we were to attempt to try all possible orderings of points, this step would be computationally intractable, with all possible combinations constituting $S!$ permutations where S is the size of the data chunk. Instead, we can randomly shuffle the points. This approach is called the Iterative Control Method, and its convergence is discussed by Julian Besang [4]. This will not guarantee that the cluster converges to the global minimum for the objective function for MCI K-means, but the algorithm will converge in relatively few iterations. Masud et al. report that where $S = 1000$, the EM algorithm converges in roughly 10 iterations, and the growth of the iteration count is sub-linear with respect to the growth of S .

The Minimisation step is to simply recalculate the centroid of each cluster, μ_i , by averaging all the points now in that cluster. The calculation for that step is:

$$\mu_i = \frac{\sum_{x \in X_i} x}{|X_i|}$$

We then repeat these steps until the convergence condition is met. We con-

sider the algorithm to have converged if no single point has changed cluster after the point re-assignment phase, as per the original specification of ICM [4].

The Micro-Clustering Step

Once the macro-clustering step is complete, we will have K macro-clusters which may or may not be class pure. Our next step is to break these macro-clusters into micro-clusters which are class pure. The procedure is rather simple based on what is considered pure:

1. Firstly, if a macro-cluster only contains labelled points from a single class, then it is already pure, and we can consider that macro-cluster as a micro-cluster.
2. Secondly, if a macro-cluster contains labelled points from two different classes, then we split that macro-cluster into two micro-clusters where each micro-cluster will only contain points from one class.
3. Thirdly, if a macro-cluster contains labelled points from one class as well as unlabelled points, we will split these points into two micro-clusters, one with the labelled points, and with unlabelled points.
4. Fourthly, if we have a micro-cluster which only contains unlabelled points, we will further subdivide that micro-cluster into micro-clusters based on the predicted class labels of those unlabelled points.

The reason we split these macro-clusters into micro-clusters is so that we can make definitive statements about the labels of the labelled micro-clusters. If this is the goal of the splitting, we might ask why we split unlabelled micro-clusters based on their predicted classes. If an unlabelled micro-cluster is constituted of two underlying classes, which have some geometric division, hopefully our predictions from previous models will be able to detect this. If we can successfully split this micro-cluster, the quality of our label propagation will increase greatly. An example is given in figure 3 taken from Woolam et al. [1]. Once we have finally decomposed our macro-clusters into micro-clusters, we will not store the complete micro-cluster, because we have no use for the complete listing of points in the micro-cluster. In the Label Stream algorithm, three basic pieces of information are stored about the micro-cluster. However, at this point, we will calculate a measure of cluster quality based on the LoOP measure, explained later:

1. The centroid of the micro-cluster, μ




Fig. 1. Illustrating micro-cluster creation. ‘+’ and ‘-’ represent labeled data points and ‘x’ represents unlabeled data points. (a) Macro-clusters created using the constraint-based clustering. (b) Macro-clusters are split into micro-clusters. (c) Unlabeled micro-clusters are further split based on the predicted labels of the data points.

Figure 3: An example of how to split macro-clusters into micro-clusters.

2. The weight of the micro-cluster, which is simply the number of points it contains, ρ
3. The label of the micro-cluster θ . Some clusters will not yet have a label.
4. In addition, we will store the LoOP score of the cluster.

This summary of the micro-cluster is called a pseudo-point, however I will continue to refer to them as micro-clusters for clarity. However, at the moment, all the micro-clusters which had no labelled points in them have no label! Hence, we move on to the third phase of our label propagation.

Label Propagation

In order to label our unlabelled micro-clusters, we will use a label propagation technique originally devised by Zhou et al[5]. Their basic technique is to build a complete graph where each micro-cluster is and each edge is a measure of weight or influence that one micro-cluster has on another. The influence in that technique that i exhibits on j is based on the Gaussian Kernel Function:

$$W_{i,j} = e^{-\left(\frac{\|\mu_i - \mu_j\|}{2\sigma^2}\right)}$$

Where μ_i and μ_j are the centroids of micro-clusters i and j , respectively, and σ is standard deviation.

To modify this affinity matrix, as it is called by Zhou, to suit our purposes, we simply need to change this undirected complete graph into a directed complete graph where $W_{i,j}$ represents the weight or influence that j exerts on i . The reason the influences will not be symmetric is that the influence j exhibits on i will now be proportional to the weight of the micro-cluster, i.e. the number of points in the cluster. This equation can be expressed as:

$$W_{i,j} = \begin{cases} e^{-\left(\frac{\|\mu_i - \mu_j\|}{2\sigma^2}\right)} \rho_j & : i \neq j \\ 0 & : i = j \end{cases}$$

We use this function to generate the affinity matrix. Note that if i and j are equal, we set the affinity matrix to be zero. When we construct our affinity matrix, we include in addition to the micro-clusters generated in this mode, labelled micro-clusters from a number of the most recent models built, so that we have a large amount of labelled data to use in our label propagation. For our experiments, we have set the number of previous models to contribute labelled micro-clusters to be 3.

We now need to normalise this affinity matrix. We will do this by taking the Normalised Laplacian of the affinity matrix. The equation for this normalisation is:

$$\mathcal{L} = D^{-1/2} L D^{-1/2}$$

Where L is the affinity matrix we built before, D is a matrix analogous to the degree matrix in a discrete setting. D is a diagonal matrix where the diagonal entries are the sum of the rows of the affinity matrix, and zero otherwise. This can be expressed as:

$$D_{i,j} = \begin{cases} \sum_x L_{ix} & : i = j \\ 0 & : i \neq j \end{cases}$$

Since this matrix is diagonal, taking the inverse and square rooting the matrix is as simple as taking the inverse of each diagonal element and then taking the square root.

Before we apply the label propagation, we need to construct a representation of the labels of our micro-clusters. We will call this matrix Υ . We will generate a C by M matrix, where M is the number of micro-clusters. For micro-cluster m , if it is unlabelled, the m th row will contain only zeros. If m is labelled, and of class c , then this matrix will have a 1 at Υ_{mc} .

Let this initial state of Υ be $\Upsilon^{(0)}$. Now we can let the labels of the Υ matrix propagate through our complete directed graph using the affinity matrix. The equation with which this happens is:

$$\Upsilon^{(n+1)} = \alpha \mathcal{L} \Upsilon^{(n)} + (1 - \alpha) \Upsilon^{(0)}$$

Here, α is the learning rate, which we set of 0.99 in our experiments as suggested by Masud et al.[3]. The label of a given micro-cluster m at convergence is the column at which the value of Υ_{mc} is greatest. The convergence condition on the label propagation is the point where no micro-cluster changes its label over subsequent iterations.

Using a Model to Classify a Point

In order to classify a given point against a single model, we build a vector Y with C elements, where Y_c represents a score for that point being of class c . To generate this score, we simply use the Gaussian Kernel Function to score the distance between the point and every micro-cluster, and then add that score to Y_c where c is the label of that given micro-cluster. We will also, after comparing the point to every micro-cluster, normalise the vector, with a small non-zero epsilon to avoid dividing by zero. This process can be expressed as an equation for labelling point x :

$$Y = \frac{\sum_j (W(x, \mu_j)) \theta_j}{\sum_j W(x, \mu_j) + \epsilon}$$

where

$$W(i, j) = e^{-\left(\frac{\|x_i - x_j\|}{2\sigma^2}\right)} \rho_j$$

and θ_j is the label of micro-cluster

To classify the point against this model, we simply label the point as class c where c is the index with the largest value in Y .

Classification and the Ensemble

An ensemble is a collection of models which we use to classify points in our stream. We use an ensemble to classify points in the data stream rather than the most recent model for several reasons. Firstly, because a single model is

trained on only a small subset of the dataset, a certain model can be poorly trained. In this case, our poorly trained model will be quickly evicted from our ensemble.

Secondly, the ensemble is an excellent way to handle concept drift within a stream. If there is a concept shift within the stream, models trained after the concept shift will gradually replace those models trained beforehand, and slowly the classification decisions of the ensemble will adapt to the concept drift. If however, a certain change in distribution of data in a chunk is caused by random chance and is not genuine concept drift, the model trained with that data will simply be evicted as soon as the data distribution returns to normal.

Adding a Model to the Ensemble

Once we have generated a model, we will need to include it in our ensemble of Models, which we will use to classify points. There are, however, certain steps we need to execute whenever we add a model to an ensemble. Firstly, we need to add our new model to the current ensemble. If our ensemble is not full, then we simply add it without any modification. However, if our ensemble is full, we need to evict our worst performing model. We do this by taking the labelled points in the data chunk used to train the most recent model (i.e., the most recent labelled points) and scoring the performance of each model in the ensemble. If two models perform equally well, we remove the oldest model.

Handling a New Class in the Data Stream

As we process the data stream, we may have a new class appear in the data stream that has never been seen before. This could be because our manual classifier felt that an examined data point could not be described within the current class framework. Woolam et al.[1] call this an evolved class. On the other hand, this could be a very rare class which has always existed in the data stream, but by chance has only now appeared. This raises the issue of what should be done about this new class.

Since previous models do not know about this new class, it is not possible that previous models can predict it. So what we will do is inject a micro-cluster representing that point into the previous models, giving them a chance to predict the point when they are being evaluated for quality when a new

ensemble is added. So, for any class which has never been seen before in the stream, we will add a micro-cluster to each model in the ensemble of that class.

However, doing this repeatedly would inflate the number of micro-clusters in each model, and possibly inflate the complexity many of our processes. So to compensate, for every injected micro-cluster, we will find the two micro-clusters in each model which are of the same label and closest together, and merge them. This merging consists of removing both micro-clusters and creating a new micro-cluster where the centroid is the average of the old centroids, and the number of points is the sum of the number of points in both clusters. If we have so many classes that each micro-cluster in a model has a distinct label (and hence, no point can be merged), then we have violated the rule of thumb that the number of macro-clusters K should be greater than the number of classes in the data stream.

Hence, we merge the nearest two micro-clusters only if:

1. There exists two micro-cluster of the same label
2. The total number of micro-clusters, $M \geq K$, the number of macro-clusters.

There is the issue of whether injecting models with the new class and merging the closest pair of micro-clusters of the same label increases the classification accuracy. On the one hand, we might think the accuracy increases because no model without information on the class can predict a point of that class. However, injecting a micro-cluster into all of the models of the ensemble could increase the correlation between each model, reducing the overall accuracy of the ensemble. [6]

The theory behind this injection is given a full discussion in Masud et al. [3], and I will not repeat it here. The key finding is that as long as single model error decreases monotonically as the prior probability of the evolved class increases, our injection will always improve the quality of the ensemble classification given that our ensemble is not too large. This conclusion is drawn by proof incorporating a proof included in Turner et al.[6], and the practical conclusion drawn by Masud et al.[3] is that the ensemble size should never exceed 10 models.

Removing Harmful Micro-Clusters from Models

When updating our ensemble, we also attempt to remove micro-clusters from models which may have a negative impact upon the classification accuracy of our ensemble.

This may seem like a tricky task, however a very clear theory for dealing with this problem is presented by Woolam et al. [1]. Firstly, we introduce the notion of a swing voter. A swing voter is a micro-cluster which determines which class a point to be predicted x will be classified as: such that if that micro-cluster were a different label, then that point would be classified differently. It is asserted that such a micro-cluster must exist, and usually, it will be the micro-cluster which has the greatest value on the inductive label propagation function, $W(i, j)$.

Lets say that the probability a micro-cluster m_i is a swing voter for a test point is α_i . So for instance if $\alpha_i = 0.5$ that means that if there are 100 test points, then m_i is the swing voter for 50 of them. Also, suppose we know the real class of a certain test point. Let p_i be the probability that for micro-cluster m_i and the test points for which it is the swing voter, the test point has the same label as m_i . So if p_i is 0.5, then for the 50 test points for which m_i is a swing voter, only 25 of them have the same label.

Using this framework, we can say that a point x is not misclassified because of a certain micro-cluster m_i if:

1. m_i is not a swing voter for that point, or
2. m_i is a swing voter and the class of x and m_i are the same.

The probability of the first case is $(1 - \alpha_i)$ and the probability of the second is $(\alpha_i(1 - p_i))$. Hence, the probability of a certain micro-cluster ensuring a point is mislabelled is $1 - \alpha_i p_i$. This means the probability that the next N test cases are not mislabelled due to m_i is $(1 - \alpha_i p_i)^N$, and the chances of a single point being mislabelled due to m_i is $1 - (1 - \alpha_i p_i)^N$. The expected classification error for a whole model is: $(1 - \sum a_i (1 - \alpha_i p_i)^N)$.

This means that if we want to reduce the classification error for the whole model, we can either reduce α_i or reduce p_i . In fact, it is this reasoning that led us to split unlabelled micro-clusters in the clustering step: by splitting the micro-cluster, we essentially replace one cluster with high p_i with two with low p_i .

However, if we were to attempt to reduce α_i , we could do that by removing

all micro-clusters! But that would be preposterous. Instead, what we will do is keep track of micro-clusters which have high p_i , and if they become too high, we will simply remove them. We call this measure the accuracy of the cluster. This should improve the model classification accuracy, because for that micro-cluster, $\alpha_i = 0$ and so it will not cause any misclassification. However, this improvement only occurs if there is no cascading effect. If we remove a certain micro-cluster, now adjacent micro-clusters are swing voters for the points for which our deleted micro-clusters were swing voters. This may or may not lead those micro-clusters to be eligible for deletion. If we continue at this rate, then most or all of the micro-clusters could be deleted.

Unfortunately, there is no tidy theory to deal with this problem. For our purposes, there is a hard limit of micro-clusters eligible to be deleted. Woolam suggested a threshold of no more than 50% point deletion, however experimental testing suggested that a more aggressive policy could yield greater classification accuracy. For experiments in this paper, the maximal number of micro-clusters removed from model could not exceed 80%.

Using LoOP to evaluate cluster quality

The previous sections lay out the Label Stream algorithm in its entirety. However, if we are to somehow attempt to use density based cluster quality measures to increase the classification accuracy of the algorithm, we need a method of evaluating the quality of the micro-clusters we generate. If we could score each micro-cluster based on its quality, then we could weight its contribution to the classification of the point, and hopefully achieve a better classification accuracy.

Cluster evaluation methods typically use two basic measures. These are intra-cluster density, and inter-cluster separation. For example, Moulavi et al [7] recently released a paper in which they evaluate cluster quality by scoring a cluster based on these two measures. To evaluate intra-cluster density, they score a cluster as the longest edge in a minimum spanning tree of the cluster, and evaluate the inter-cluster distance by finding the nearest pair of points between two clusters. They can then build an overall score based on these two measures.

However, evaluating the quality of micro-clusters should not use inter-cluster measures. This is because cluster overlapping is not necessarily a sign of bad clustering. If we take a macro-cluster which contained labelled points

of a class, and unlabelled points which are actually of that class, then we will split that cluster into two different micro-clusters at the micro-clustering phase. However, these overlapping clusters are not a problem because they both come from the same underlying class distribution.

Another viable method to change the micro-clusters is simply to find outliers in a cluster and remove them. Since our clusters centroids are mean based, an outlier in a small cluster can skew the centroid of a cluster significantly. Identification and removal of these points maybe another way

In this section, we will discuss LoOP [13], a method which evaluates how likely it is a single point is an outlier within a cluster. After we have explained how this measure works, we will look at how it can be applied to this problem.

LoOP: The Local Outlier Probability Factor

When trying to determine whether a point is an outlier, many outlier scores only give an arbitrary score which must be interpreted. However, LoOP is innovative in the way that rather than producing an arbitrary score, each point is assigned a probability of being an outlier. This makes the results much easier to interpret. The steps for calculating the LoOP score are detailed more extensively in the work of the authors of the paper, Kriegel et al.[?], but I will recite the main points here.

The first step is to define a probabilistic distance. Suppose we have a distance measure like euclidean distance, which we denote as $d(x, y)$. Suppose also we have some point we are interested in o , and a set of objects S . Then we want our probabilistic distance between a point o and the set of objects S to be have the property:

$$\forall s \in S : P[d(o, s) \leq pdist(o, S)] \geq \phi$$

where ϕ is some constant. Intuitively, what this means is that a sphere around o with a radius $pdist(o, S)$ covers any element in S with probability ϕ . Now, if we wish to estimate probabilistic density, we can take the inverse of the $pdist$:

$$pdens(S) = \frac{1}{pdist(o, S)}$$

Now, if we use $\lambda = \sqrt{2}erf^{-1}(\phi)$ rather than ϕ , where erf is the gaussian error function, we can interpret λ as a multiplier of standard deviation σ . So

for instance, $\lambda = 1 \Leftrightarrow \phi = 68\%$, and $\lambda = 2 \Leftrightarrow \phi = 95\%$, etc. Now we can define a standard distance between one point o and a cluster of points S :

$$\sigma(o, S) = \sqrt{\frac{\sum_{s \in S} d(o, s)^2}{|S|}}$$

and hence, we can define the probabilistic distance as:

$$pdist(o, S) = \lambda \sigma(o, S)$$

We need to compare the probabilistic distances of each point in a cluster of points in order to evaluate how likely it is that one point is an outlier. In order to do this, we will firstly compare the $pdist$ of a single point to all the $pdist$ of every point in that cluster. For that purpose, we define the Probabilistic Local Outlier Factor (PLOF) to be:

$$PLOF_{\lambda, S}(o) = \frac{pdist(\lambda, o, S)}{E_{s \in S(o)}[pdist(\lambda, s, S)]} - 1$$

where E refers to the expected value. This PLOF score is not yet probabilistic, and we cannot compare the outlierness of different points in different distributions yet. So we need a measure NPLOF, which stands for Normalised Probabilistic Local Outlier Factor. This is defined:

$$NPLOF = \lambda \sqrt{E[(PLOF^2)]}$$

The purpose of this measure is to standardise what constitutes an outlier across different contexts. This could either be within a cluster, or between different clusters. Finally, our measure is still not probabilistic, so we define the LoOP score of such a point to be probabilistic using the Gaussian Error Function as below:

$$LOOP_S(o) = \max\{0, erf(\frac{PLOF_{\lambda, S}(o)}{nPLOF\sqrt{2}})\}$$

The resulting LoOP score will be close to 0 for points which are in dense clusters, and close to 1 for points that are density based outliers.

Using LoOP scores to improve Cluster Quality by removal

One way to improve the quality of a given micro-cluster is to simply remove outliers. The reason that this will improve the quality of the cluster is that when we calculate the centroid of a cluster, we calculate the position of the centroid by averaging the positions of every point in the cluster:

$$\mu_i = \frac{\sum_{x \in X_i} x}{|X_i|}$$

Hence, a single outlier has the possibility to skew the position of the centroid of a cluster. Of course, outliers can cancel each other out, but simply removing an outlier from a cluster will not make the cluster quality worse.

The other question to ask is why LoOP is more appropriate than other outlier factors like LOF[14]. The reason why we prefer LoOP over other outlier measures is that LoOP gives us a way to standardise "outlierness" between clusters. For instance, if a cluster has four very close points and one point further away, that point may be interpreted as an outlier. However, when we consider the distances between points in other clusters, it may become apparent that in the context of the dataset that one further away point is not an outlier.

The procedure for calculating the LoOP score for every point is as follows:

1. We calculate the PLOF for every point in relation to the cluster it is in.
2. We then calculate the NPLOF of the entire model by calculating the expected value of every PLOF in every cluster, as we saw in the description of LoOp.
3. Now we can calculate the LoOP score of every point.

We will remove a point given it has a LoOP score greater than 0.5, with the LoOP calculations done with $\lambda = 3$. We can interpret this as removing a point if it is more likely than not that it is outside 3 standard deviations.

Using LoOP scores to predict cluster quality

A more uncertain method of improving the classification accuracy of the algorithm is to use the LoOP scores of individual points within a cluster to attempt to score the whole cluster, and then use that score to bias the classification process towards micro-clusters that have a higher cluster quality. Provided we had an appropriate measure, we could simply weight each micro-cluster by its quality when considering it in classification. Let τ_i denote the cluster quality of micro-cluster i . Then we can change our classification influence function to be:

$$W(i, j) = e^{-\left(\frac{\|x_i - x_j\|}{2\sigma^2}\right)} \rho_j \tau_i$$

We need our cluster quality measure to be a decimal value between 0 and 1, where 0 denotes that the micro-cluster should be entirely disregarded. Fortunately, LoOP ensures that all points are given a LoOP score between 0 and 1. So what function should we use to aggregate the LoOP scores of our clusters?

The minimum and maximum functions are candidates, but not very good ones. Suppose for instance we let the LoOP score of a cluster be:

$$\tau_i = 1 - \max(s \in i)$$

If a cluster had a point that was a certain outlier, say a 99% outlier, $\tau_i = 0.01$, so that would mean we effectively ignore the entire cluster. This is not a very good idea, because the cluster may be very dense and informative aside from that one outlier. Simply put, a function based on the value of one point is probably not going to yield a good result.

Another possible candidate is the average of the LoOP scores. However, here we have a problem of dampening, where due to the fact all clusters will consist primarily of low LoOP score points, we will be unable to distinguish significantly between good and poor quality clusters. For instance, suppose a cluster is 20% heavy outliers. If these outliers are all corroborating to distort the centroid of the cluster, this could seriously damage the quality of this cluster. However, using a measure like:

$$\tau_i = 1 - \frac{\sum_{x \in X_i} x}{|X_i|}$$

will yield us $\tau_i = 0.8$, not a significant enough distinction between high and low quality clusters.

The solution to both these problems is to use an average of the points in the cluster which have the highest outlier factor. It makes sense that we can score the whole clusters quality on these few points, because it only takes a few points to entirely distort the centroid of the cluster. So, in order to determine cluster quality using LoOP score, we:

1. Find the LoOP scores of all points in each cluster.
2. We then, for each cluster i , sort the points by score.
3. We take the highest 5% of these scores, and average them $Av(i)$.
4. $\tau_i = 1 - Av(i)$.

Interaction between Micro-cluster Removal and Proposed LoOP measures

It is worth noting that the algorithm already has a built in notion of cluster quality; the accuracy of a micro-cluster. This is extensively discussed in a previous section, however the basic notion is that we assign every labelled point in a training set to its nearest micro-cluster, and if a certain micro-cluster has too many points of the wrong label to it, it is removed. Hence, this method does not attempt to modify the micro-cluster or reduce its influence somewhat, but rather entirely remove it.

The methods experimented with in this paper are different in two respects. Firstly, the methods of evaluation used with density based outlier methods evaluate quality once, independent of the current training data, rather than iteratively and dependent on the state of the current data chunk. Secondly, the above methods attempt to only dampen the influence of micro-clusters which contain outliers, rather than exclude them.

The other thing we might wonder is how these quality modifications interact with one another. Firstly, outlier removal and micro-cluster removal by accuracy could compliment one another. If outliers skew the centroid of a certain cluster, it may move the centroid of the cluster away from test points of the same label and towards those with other labels, leading to be classified as an inaccurate micro-cluster, and removed. If we remove the outliers, an inaccurate micro-cluster may become accurate and contribute to classifier accuracy rather than hinder it.

On the other hand, it is less clear how scoring cluster quality and micro-cluster removal will interact. It is possible that micro-cluster removal entirely subsumes any effect of the cluster quality scoring.

Complexity

The complexity of Label Stream is well discussed Masud et al.[3], and I will summarise the results here.

- The complexity of the initial cluster step is $O(KSg(S))$, where $g(S)$ is the sublinear function representing how long it takes for ICM to converge, K is the number of macro-clusters and S is the data chunk size.
- The complexity of micro-cluster splitting, label propagation is M^3L , where M is the number of micro-clusters and L is the number of models in the ensemble. This complexity is largely due to the matrix operations in the label propagation.
- Refining the ensemble takes $O(ML(P/100)S)$ time where P is the percentage of labelled points required before we can build a model.

Since M is bounded by $K(C + 1)$, and K , C , L and P are relatively small, we can talk about Label Stream being approximately linear with the size of the size of the data chunk, S .

The complexity of LoOP is not much different. To calculate the nPlof, we need to calculate the PLOF of every point the the data chunk. To calculate the Plof of a point, we need to calculate the pDist of every point in that microcluster. To do that, we need to compare distances with every point in the micro-cluster. So suppose that A is the average size of a micro-cluster. So:

$$S = MA$$

then the complexity of LoOP is $O(SA^2)$, which is approximately linear in the size of S , since A will be much smaller than S .

Experiments

Experiments were conducted on a variety of datasets. Since the aim of the alterations to the Label Stream algorithm was to increase the classification accuracy of the algorithm, no experiment tracked the relative point classification rate (speed) of the algorithm with and without the LoOP density score. Subsequently, the hardware the tests were performed on is not particularly relevant, however for completeness they were performed on a MacBook Pro with a 2.5GHz processor and 4GB of memory.

The experimental method is as follows. Because Label Stream was developed in comparison to an algorithm called On Demand Stream [16], the testing method follows the technique used by the ODS developers. What we do is split the data stream into two streams. The first data stream we use in the ordinary course of the algorithm, partially labelling data and breaking the stream into chunks, building models and placing them into the ensemble. The second stream we use to test the accuracy of the classification, however, we still break that stream into chunks. We will test on data chunk $i + 1$ after chunk i has been used to train a model and modify the ensemble.

Each dataset is run three times. The first is a control run, with no outlier removal based on LoOP score nor cluster score weighing based on the LoOP cluster. This is a test which replicates the original design of Label Stream [1] Then, we run the same algorithm except we remove points which we consider to be outliers using the LoOP score after the micro-clustering phase. Thirdly, we do not remove outliers, but rather use the LoOP scoring to influence the classification without modifying the cluster. Since the Label Stream algorithm is non-deterministic, for each of these three permutations of the algorithm, the test was run twenty times and the cumulative accuracy results at each data chunk was averaged.

When graphing the results, each line corresponds to one type of run on each data set. Every so many iterations, each dataset will display a standard error to indicate the variance over the different accuracies for that test type on that iteration. These appear on the graphs like error bars.

Relevant parameters:

- Number of contiguous models to take labelled micro-clusters from in label propagation = 3
- σ for label propagation and classification = 0.25
- α learning rate for label propagation = 0.99

- K , the number of macro-clusters used in the macro-clustering step = 50
- P , the percent labelled data in a chunk used to train a model = 10%
- r , the number of models we take labelled micro-clusters from to use in label propagation = 3
- S , chunk size, is 1600, except for the synthetic dataset, which is 1000.
- L , the size of the ensemble = 6.
- The accuracy threshold for a micro-cluster deletion = 0.7
- Limit of micro-cluster removal in deletion phase = 80%
- λ for LoOP scoring = 3
- Boundary LoOP score for a point to be removed in outlier removal = 0.5
- Percentage of points used in LoOP cluster quality evaluation = 5%

SynD Dataset

The Dataset

The primary purpose of the SynD dataset is to examine how each algorithm deals with conceptual drift in the data stream. SynD is a synthetic dataset generated in MOA [15] which is based on a shifting hyperplane. The equation of the hyperplane is:

$$\sum_{i=1}^d a_i x_i = a_0$$

where d stands for the number of dimensions, a_i is the weight associated with dimension i , and x_i is the value of the i th dimension of data point x . Each point can be classified by considering whether $\sum_{i=1}^d a_i x_i \leq a_0$; if it is, then the point is placed in class 1, otherwise it is placed in class 2. Each point is generated vector of length $d\{x_1, \dots, x_d\}$ where $x_i \in [0, 1]$. The weight vector is also a vector of length $d\{a_1, \dots, a_d\}$ where each weight is in the range of $[0, 1]$. The value of a_0 is set to generate roughly the same number of points from each class. We can do this easily if we set a_0 as the average of the

weights, $\frac{1}{2} \sum_{i=1}^d a_i$. Also, we swap the class label on some points to generate noise in the dataset.

We can introduce concept drift by choosing a subset of the weights and gradually varying them over a certain number of data points in the stream. We simply define a magnitude of drift t , a vector of drift direction $\{s_1, \dots, s_d\}$, and over N data points a certain subset of the weights so that a_x changes by $s_x t / N$. Finally, after N data points have been generated, we can generate a new subset of weights, a new magnitude of change, and randomly change the direction of change of weight for some weights.

For our experiments, we followed the experimental parameters chosen by the developers of the LabelStream algorithm for comparative reasons. The parameters are:

- Instances = 250,000
- $d = 20$
- Number of drifting attributes = $0.2d$
- Magnitude of drift $\in [0.1, 1]$
- Chance of drift direction change = 10%
- Number of points before drift recalculation = 1000
- Noisy points = 5%

Results

As we can see from Figure 4, using cluster weighting or outlier removal had little effect on the cumulative accuracy of the algorithm on this dataset. Due to the low class count relative to the number of macro-clusters, the cluster quality is a lot better than some datasets with more classes.

We also note that for the synthetic dataset, there was a great deal of variance in the results, as shown in the standard error bars for each kind of run. This is due partially to the fact that only 10% of the data is labelled, and partially




Figure 4: Cumulative accuracy for Data set with different quality settings

due to the randomness in the way the concept drift was generated in the synthetic data generator. This variance casts further doubt on whether the results in the graph represent any serious distinction between the methods.

KDDCup Network Intrusion Dataset

The Dataset

The KDDCup Network Intrusion dataset is an extremely common test dataset for Machine Learning problems. Each data point contains 41 categorical or real value attributes relating to the status of a network connection, and then a classification value which states whether the network connection was an attack, and if so, what kind of attack it was. The classification problem is hence to identify whether the connection is normal or whether it is an attack of a certain variety. However, for my implementation, all categorical data except the class of network attack was removed for purposes of comparison with the Label Stream algorithm. The actual set used was the fully labelled ten percent subset freely available from most data repositories ¹ which has approximately 500,000 data points.

It is worth noting that this data was not originally a stream, it was simply interpreted as a stream. This means that extremely sudden concept drift can occur in the data which results in large changes to the ensemble. It is also worth noting that although there are 24 classes in the dataset, 98% of the data points fall into one of three classes: Normal, Neptune attack, or Smurf attack. This means the success of the classification depends largely on predicting these three classes.

Results

The KDD dataset seems to be one of the few datasets which suggests that outlier removal may be a feasible technique. There is, at points, several cumulative percentage points between the outlier removal method and the other two methods at various points in the stream. However, the KDD dataset is another dataset which exhibits a great amount of variance in the results. The great spikiness of the cumulative accuracy for these graphs is due to the fact that the KDD dataset often changes very abruptly. So for instance, you might have trained all the models in your ensemble to handle

¹www.kdd.ics.uci.edu/databases/kddcup99/kddcup99.html




Figure 5: Control, Outlier removal and cluster scoring compared on the KDDCup dataset.

normal network traffic, then a concerted 10 model block of new network attacks could appear, and then return to normal. So KDD is a dataset with abrupt concept drift. However there is no source of extra randomness in the data except for which points are labelled unlike the Synthetic dataset. So while this data set may appear to suggest the outlier removal approach shows promise, we need to be extremely weary in the context of the great variance of the results.

Forest Cover Dataset

The Dataset

The Forest Cover Dataset is a dataset which attempt to classify what type of forest cover is in a particular area of forest based on certain attributes. Each data point has 54 attributes, however most of these are binary variables relating to particular characteristics of soil type, so if you count them as one attribute, there are only 10 attributes. All of these attributes are integers or real values. There are then 7 types of forest cover into which the data point can be classified.

The dataset is large, containing 500,00 data points. These are predominantly made up of two classes: Spruce-Fir with 210,000 records and Lodgepole Pine with 280,000 records. The remaining classes number between 2,500 and 40,000 records. It is also worth noting this dataset is not a stream, it is simply interpreted as a stream.

Results

The forest cover data set sets out the strongest case for the outlier removal technique being effective. Firstly, there appears to be separation of perhaps a single percentage between the outlier removal towards the end of the stream. We need to remember that this is a cumulative accuracy measure, so a single percent can be quite significant. Secondly, the variance which plagues the previous two data sets does not appear to be present in this data set, despite the fact it was still working with only 10% labelled data. This is very promising, however the cluster removal technique still looks to be ineffective.




Figure 6: Cumulative accuracy for Data set with different quality settings

Power Usage Dataset

The Dataset

The Power Usage dataset is a very simple data stream. There are only two attributes for each data point; the amount of power on the main grid and the amount of power transformed from other grids. The class to be predicted is the hour of the day which these power measurements came from. This means there are 24 classes. The interesting property of this dataset is that many of the classes are extremely similar. In fact, when this data stream was used for experimentation by Zhang et al.[2] they transferred the data set into a binary classification problem, possibly to avoid this problem. We also followed a similar approach. Here there was a distinct difference in data between the hours that people are sleeping (1am - 7am) and all other times, so we divided the data into a binary classification problem based on that.

The other interesting aspect of this dataset is that there is natural concept drift, caused by factors like the day of the week the data is measured on and the season. The downside to this dataset there are only 30,000 records, which is very small. What we did was rerun the data over and over until there was a great enough size to do regular testing on.

Two Class Results

Here we see something entirely different to previous datasets; An extremely small amount of variance relative to the other data sets, and very little difference between any of the methods. What this suggests is that the clustering was of sufficiently high quality that barely any outliers were removed, and no clusters were penalized for having a low cluster quality score. This seems entirely possible considering there are only two classes and 50 macro-clusters to build clusters out of. This is also the case for the Synthetic dataset, but there results were much different. One of the possible reasons for this result is that the data chunk size was too large to encapsulate any of the seasonal concept drift in the model. For instance, we use a 1600 point data chunk to train models and another 1600 points to test. That amounts to 3200 points per model. There is one data point per hour, and there are 24 hours in a day, which means one model encapsulates 133 days, and 6 models encapsulate two years of data. This would mean that there would be concept drift occurring within models rather than between them, and so between iterations the data would seem reasonably static. This would make the data extremely messy.

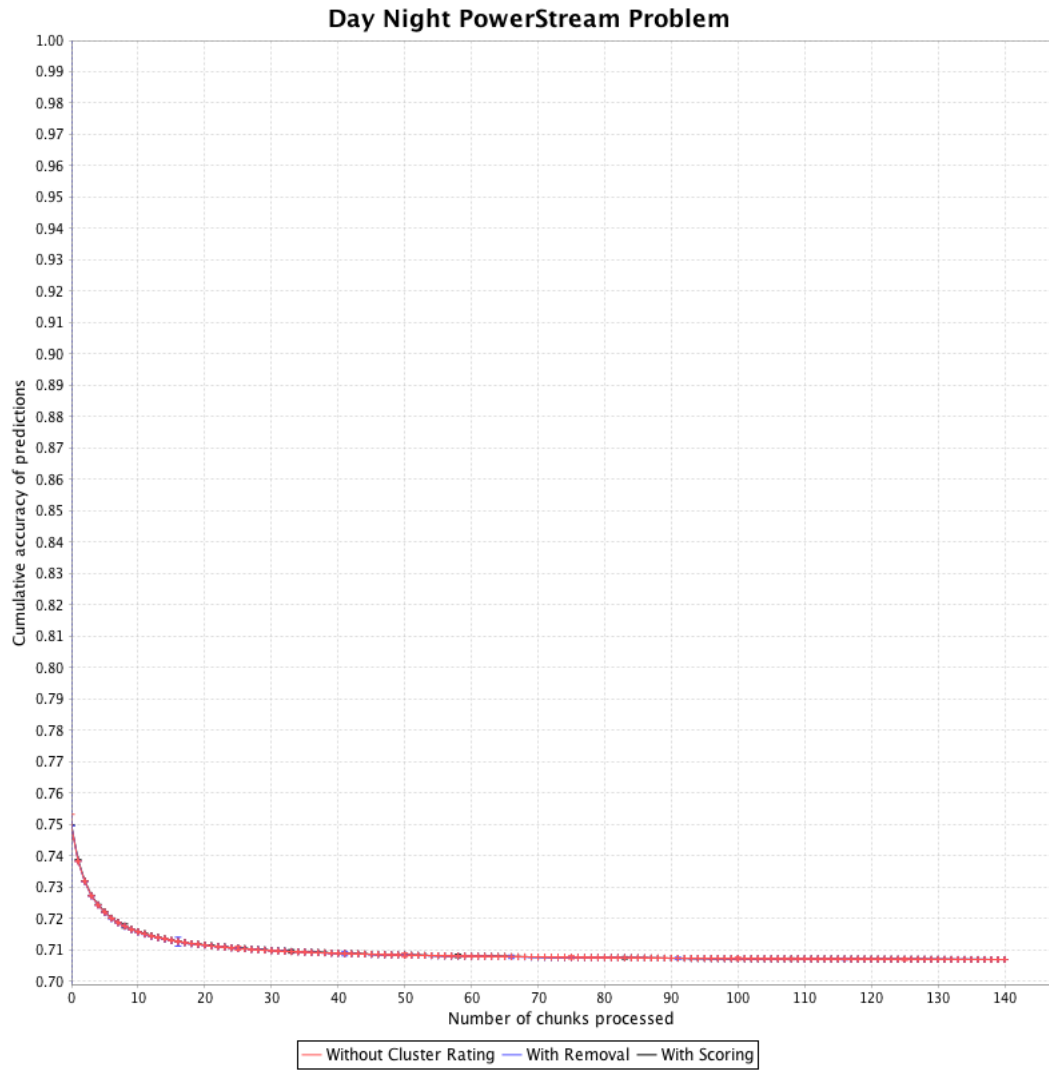


Figure 7: Tests on the Binary Power Stream Data Set

So given the size of the chunks and the number of data points in the data set, this un-interesting result seems explicable.

Sensor Stream Dataset

The Dataset

The sensor stream dataset is a real world data stream generated by the Intel Berkeley Research Lab. They planted sensors throughout their offices which every few minutes took measurements of temperature, humidity, light and voltage. The classification problem is to decide which of the 54 sensors the data point came from. One of the interesting features of this data is the sheer volume, since it has over 2,000,000 records. However, we must consider that per class, the amount of data is probably in line with some of the other data sets we have tested here.

Results

Results Overall

There seem to be some interesting trends in the tests. Firstly, no method significantly out-performed another over any dataset. There were minor variations, and certainly cumulative accuracy can make substantial differences look small, but as a point of comparison, the decision on whether to consider labelled or unlabelled points as a distinct class made more of a difference to the Synthetic dataset than any of these methods. Certainly, the cluster score technique did not perform any differently to the control technique.

Secondly, the outlier removal technique did perform better than the default method on two data sets. What is most fascinating, however, is that these two data sets were the two data sets which were not designed to be stream data. This of course could be purely by chance due to the high variances of the KDD and SynD data sets. However, the idea that outlier removal is most helpful on the datasets which have no inherent structure is one that seems intuitive. If a file is not a genuine data stream, an outlier cannot be indicative of an incoming trend. However, if the data is a genuine data stream, an outlier can either be a sign of bad clustering, a simply outlying point, or indicative of an incoming concept drift. By removing outliers, we

perhaps are inhibiting the detection of genuine concept drift, which SynD certainly contained, unlike the KDD and Forest Cover datasets.

In addition, we will perform a Mann Whitney U-test to examine the rankings of the various methods on each dataset [12]. To summarise, here is the table. Where results were too close to distinguish, we simply give all methods the median rank of 2.

| Data Set | Without Cluster Rating | With Removal | With Cluster Scoring |
|---------------|------------------------|--------------|----------------------|
| SynD | 1 | 3 | 2 |
| KDDCup | 3 | 1 | 2 |
| Forest Cover | 3 | 1 | 2 |
| 2 Class Power | 2 | 2 | 2 |
| 4 Class Power | 2 | 2 | 2 |
| Sensor Stream | 2 | 2 | 2 |
| Total | 13 | 11 | 12 |

For a comparison between the control and the outlier, we need to assign a Mann-Whitney rank score:

- Ranking of 1: 2 pts
- Ranking of 2: 6 pts
- Ranking of 3: 11pts

This is sorting the ranks in the pair of tests we want to compare, and assigning each score its position in a sorted list. If items are equal, they get the average of their positions. So for instance, Rank 1 gets a score of 2 because $(1 + 2 + 3)/3 = 2$.

Now the control has a corresponding score of 42, and outlier removal has score of 33. The larger of these two, is tx . The formula for computing the Mann-Whitney number U is:

$U = n1.n2 + nx((nx + 1)/2) - tx$ where $n1$ is the number of sample the first group came from (here 6) and $n2$ is the number of sample $n2$ came from (also 6). nx is the n that corresponds to tx . Hence, we can compute the U value:

$$U = 6x6 + 6(7/2) - 42 = 15$$

Looking up a Mann-Whitney value table, the value for a 6,6 test to be less than 0.05 significance level is $U = 5$. This means we can consider this ranking of techniques highly unlikely to occur by chance alone if $U \leq 5$. Hence, we

can conclude that the ranking of the control and the outlier is very likely to be caused by random chance alone, rather than some trend across the datasets.

Similarly, if we compare the control and the cluster scoring technique:

- Ranking of 1: 1 pt
- Ranking of 2: 6 pts
- Ranking of 3: 11.5 pts

Here the score for the control is 41 and the score for the cluster ranking is 36. So:

$$U = 6 \times 6 + 6(7/2) - 41 = 15$$

So again, this pattern of rankings is also likely created by chance.

This statistical analysis seems to agree with what we can intuit from looking at each individual test case: Across the different dataset, there seems to be no clear trend emerging as to whether the control method or the outlier removal or cluster scoring technique are better, worse, or even different.

Limitations and Future Work

This study has various limitations. Firstly, we barely touched on the performance aspect of the modifications to the Label Stream algorithm. We have looked at the complexity of the LoOP calculation, but one of the strengths of the micro-clustering approach of Label Stream is that as soon as the label propagation is over, all the points in the micro-cluster can be thrown away and only the summary of the point remains. This is still the case, however calculating the LoOP score of every point in every cluster may change the memory profile of the algorithm and change its memory consumption, as well as slightly hurt the classification rate.

Secondly, a glaring omission from this report is an examination of the parameter sensitivity of various techniques. The outlier removal method, cluster scoring method, and control method involving micro-cluster removal all are sensitive to certain parameters, and the efficacy of certain techniques are almost certainly dependent on this sort of experimentation. This would also involve greater experimentation with various data sets, particularly synthetic

streams. This is the most pressing area which needs further work in the future, and will certainly yield the most insight.

Beyond that, our study was artificially restricted to only consider the LoOP score as a measure of cluster quality. In the further future, more investigation into density based clustering evaluators would be appropriate, Beyond that, non-density based cluster measures could be considered, although they may only be appropriate on data we know is distributed in a certain way.

Conclusion

Bibliography

- [1] Woolam, Clay, Mohammad M. Masud, and Latifur Khan. "Lacking labels in the stream: classifying evolving stream data with few labels." *Foundations of Intelligent Systems*. Springer Berlin Heidelberg, 2009. 552-562.
- [2] Zhang, Peng, Xingquan Zhu, and Li Guo. "Mining data streams with labeled and unlabeled training examples." *Data Mining, 2009. ICDM'09. Ninth IEEE International Conference on*. IEEE, 2009.
- [3] Masud, Mohammad M., et al. "Facing the reality of data stream classification: coping with scarcity of labeled data." *Knowledge and information systems* 33.1 (2012): 213-244.
- [4] Besag, Julian. "On the statistical analysis of dirty pictures." *Journal of the Royal Statistical Society. Series B (Methodological)* (1986): 259-302.
- [5] Zhou, Dengyong, et al. "Learning with local and global consistency." *Advances in neural information processing systems* 16.16 (2004): 321-328.
- [6] Tumer, Kagan, and Joydeep Ghosh. "Error correlation and error reduction in ensemble classifiers." *Connection science* 8.3-4 (1996): 385-404.
- [7] Moulavi, Davoud, et al. "Density-based clustering validation." *Proceedings of the 14th SIAM International Conference on Data Mining (SDM)*, Philadelphia, PA. 2014.
- [8] Hulten, Geoff, Laurie Spencer, and Pedro Domingos. "Mining time-changing data streams." *Proceedings of the seventh ACM SIGKDD international conference on Knowledge discovery and data mining*. ACM, 2001.

- [9] Domingos, Pedro, and Geoff Hulten. "Mining high-speed data streams." Proceedings of the sixth ACM SIGKDD international conference on Knowledge discovery and data mining. ACM, 2000.
- [10] Wang, Haixun, et al. "Mining concept-drifting data streams using ensemble classifiers." Proceedings of the ninth ACM SIGKDD international conference on Knowledge discovery and data mining. ACM, 2003.
- [11] Li, Peipei, Xindong Wu, and Xuegang Hu. "Mining recurring concept drifts with limited labeled streaming data." ACM Transactions on Intelligent Systems and Technology (TIST) 3.2 (2012): 29.
- [12] Siegel, S. "The Mann-Whitney U test." Nonparametric Statistics for the Behavioral Sciences (1956): 116-127.
- [13] Kriegel, Hans-Peter, et al. "LoOP: local outlier probabilities." Proceedings of the 18th ACM conference on Information and knowledge management. ACM, 2009.
- [14] Breunig, Markus M., et al. "LOF: identifying density-based local outliers." ACM Sigmod Record. Vol. 29. No. 2. ACM, 2000.
- [15] Bifet, Albert, et al. "Moa: Massive online analysis." The Journal of Machine Learning Research 11 (2010): 1601-1604.
- [16] Aggarwal, Charu C., et al. "A framework for on-demand classification of evolving data streams." Knowledge and Data Engineering, IEEE Transactions on 18.5 (2006): 577-589.